

# Spring Batch Introduction

14th Feb 2013

**Petr Zapletal**

Common Application Development

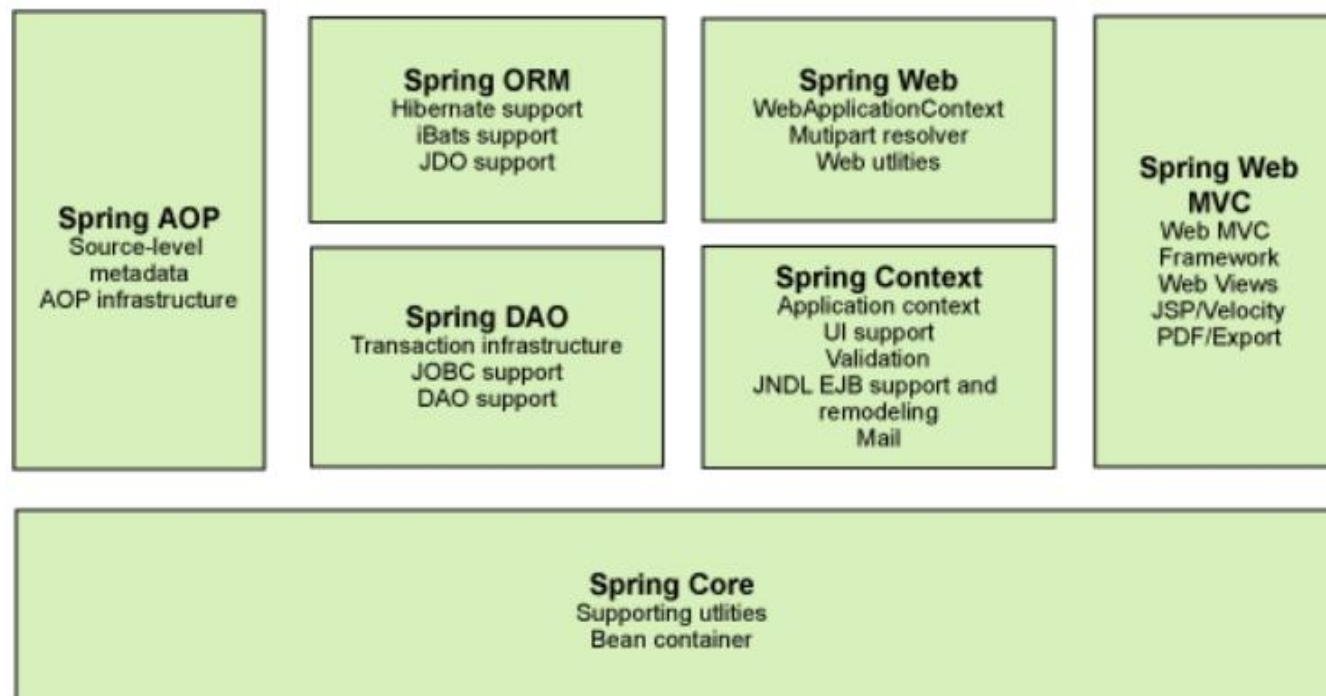


# Agenda

- Spring & Spring Batch Overview
- Domain Language of Batch
- Configuring and Running a Job
- Advanced Features
- Summary

# Spring Overview

- One of the most popular open-source application development framework for enterprise Java
- Provides lightweight container



# Spring Batch Introduction

- Java based framework for batch processing
- A lightweight, comprehensive batch framework
- Builds upon the productivity, POJO-based development approach, known from the Spring Framework
- Current version is 2.1.9.RELEASE

# Batch Processing

- What is a Batch Application?
  - Batch applications need to process high volume business critical transactional data
  - A typical batch program generally
    1. reads a large number of records from a database, file, or queue
    2. processes the data in some fashion, and
    3. then writes back data in a modified form

# Usage Scenarios

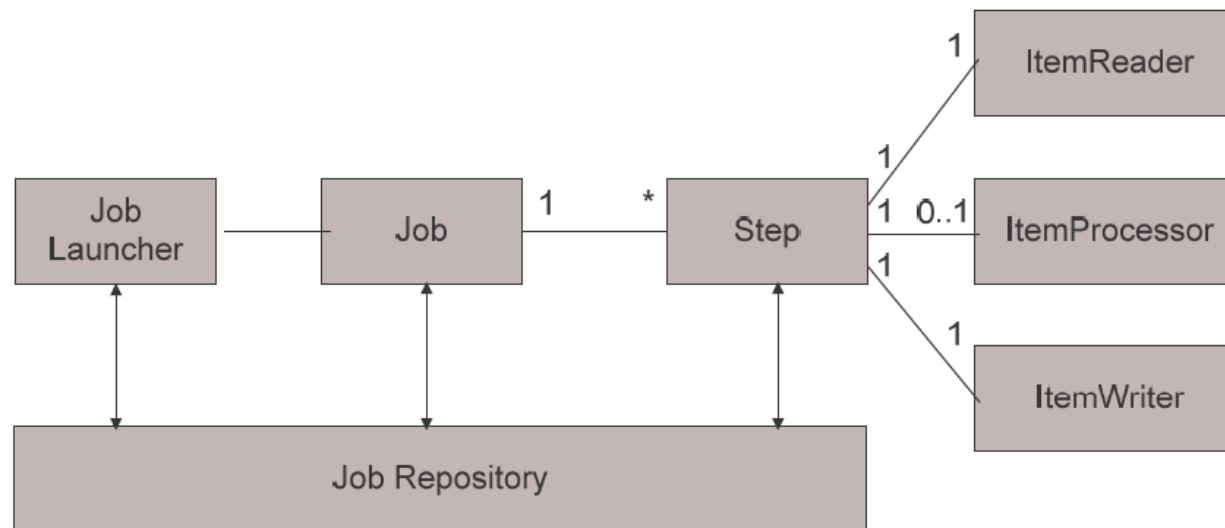
- Commit batch process periodically
- Concurrent batch processing: parallel processing of a job
- Staged, enterprise message-driven processing
- Massively parallel batch processing
- Manual or scheduled restart after failure
- Sequential processing of dependent steps
- Partial processing: skip records
- Whole-batch transaction: for cases with a small batch size or existing stored procedures/scripts

# Agenda

- Spring & Spring Batch Overview
- Domain Language of Batch
- Configuring and Running a Job
- Advanced Features
- Summary

# Domain Language of Batch I

- A job has one to many steps
- A step has exactly one ItemReader, ItemWriter and optionally and ItemProcessor
- A job needs to be launched (JobLauncher)
- Meta data about the running process needs to be stored (JobRepository)






# Domain Language of Batch II

- Job encapsulates an entire batch process
- Job Instance
  - Refers to the concept of a logical job run
  - Job running once at end of day, will have one logical JobInstance
  - Each JobInstance can have multiple executions
- Job Execution
  - Refers to the technical concept of a single attempt to run a Job
  - An execution may end in failure or success, but the JobInstance will not be considered complete unless the execution completes successfully
- Job Parameters
  - Is a set of parameters used to start a batch job
  - JobInstance = Job + JobParameters

# Domain Language of Batch III

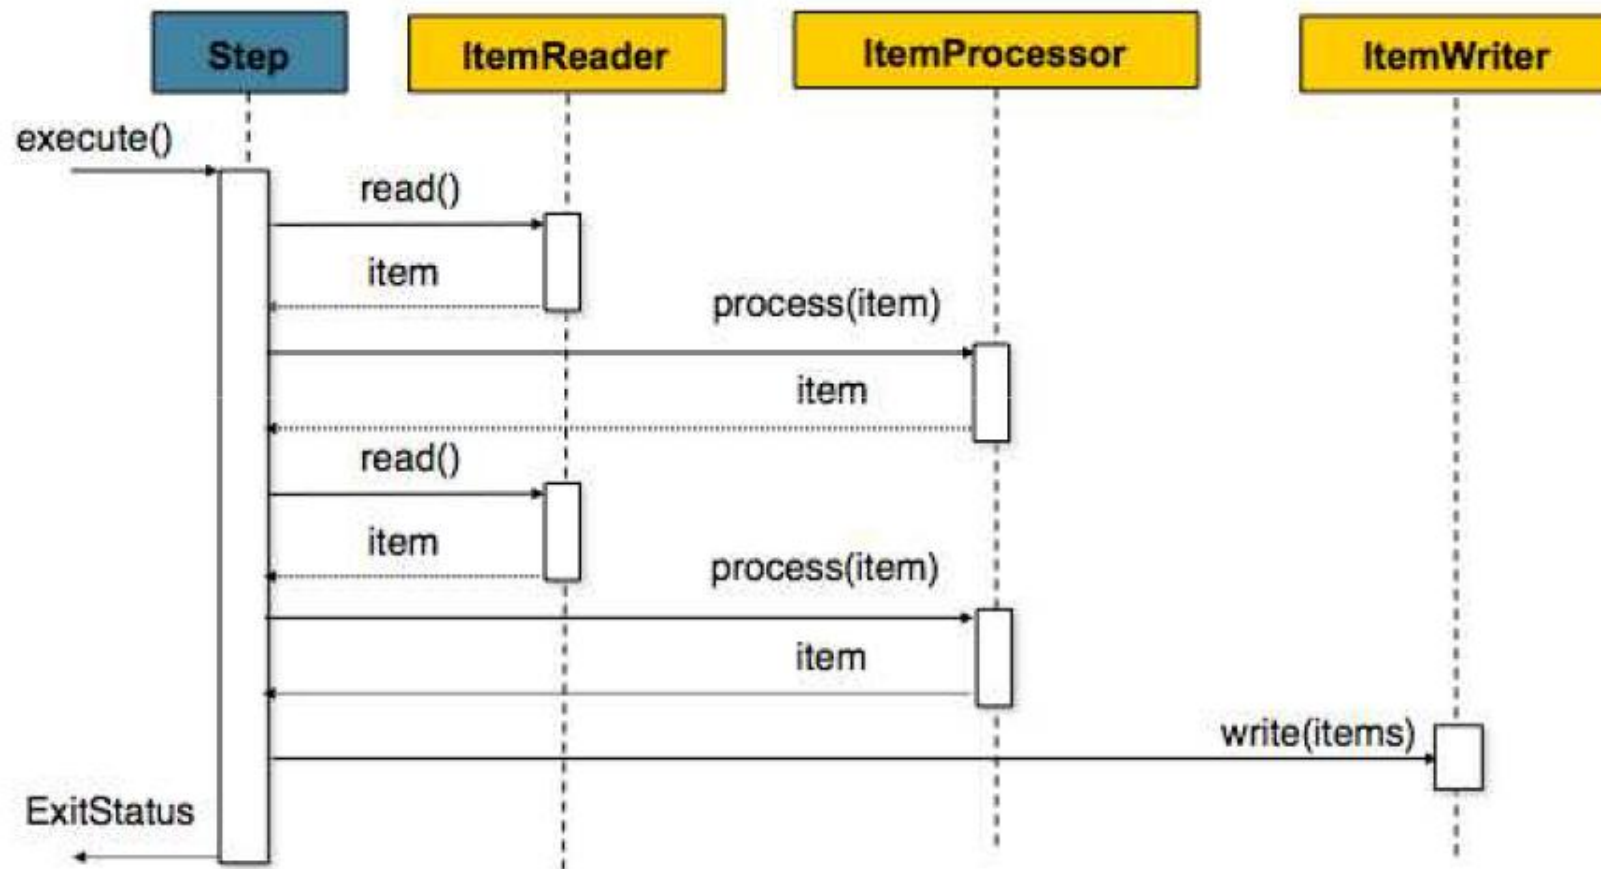
- Step
  - A domain object that encapsulates an independent, sequential phase of a batch job
  - Can be as simple or complex as the developer desires
- Step Execution
  - Represents a single attempt to execute a Step
  - A new StepExecution will be created each time a Step is run, similar to JobExecution
  - A StepExecution will only be created when its Step is actually started



# Domain Language of Batch IV

- Item Reader
  - An abstraction that represents the retrieval of input for a Step
  - When it has exhausted the items it can provide, it will indicate this by returning null
  - Various implementation available
- Item Writer
  - An abstraction that represents the output of a Step
  - Generally, an item writer has no knowledge of the input it will receive next
  - Various implementation available
- Item Processor
  - An abstraction that represents the business processing of an item
  - Provides access to transform or apply other business processing

# Readers & Writers & Processors



# Readers & Writers & Processors

## – ItemReader

```
public interface ItemReader<T> {  
    T read() throws Exception, UnexpectedInputException,  
        ParseException;  
}
```

## – ItemWriter

```
public interface ItemWriter<T> {  
    void write(List<? extends T> items) throws Exception;  
}
```

## – ItemProcessor

```
public interface ItemProcessor<I, O> {  
    O process(I item) throws Exception;  
}
```

# Job Repository & Job Launcher

- Job Repository
  - Batch persistence mechanism
  - Database and In-Memory implementation
  - Data about job executions, start times, durations, results, ...
- Job Launcher
  - Represents a simple interface for launching a Job with a given set of JobParameters

# Agenda

- Spring & Spring Batch Overview
- Domain Language of Batch
- **Configuring and Running a Job**
- Advanced Features
- Summary

# Job Definition

- Configuring a Job and its steps

```
<job id="sampleJob">
  <step id="step1" job-repository="jobRepository"
        transaction-manager="transactionManager">
    <tasklet reader="itemReader" writer="itemWriter"
             commit-interval="10"/>
  </step>
</job>

<bean id="itemReader" ...>
<bean id="itemWriter" ...>
```



# Job Repository & Job Launcher Definition

## – Configuring a Job Repository

```
<job-repository id="jobRepository"  
    dataSource="dataSource"  
    transactionManager="transactionManager"  
    isolation-level-for-create="serializable"  
    table-prefix="BATCH_" />
```

## – Configuring a Job Launcher

```
<bean id="jobLauncher"  
    class="...batch.execution.launch.SimpleJobLauncher">  
    <property name="jobRepository" ref="jobRepository" />  
</bean>
```

# Demo

```
<bean id="propertyPlaceholderConfigurer"
      class="org.springframework.beans.factory.config.PropertyPlaceholderConfigurer">
  <property name="systemPropertiesModeName" value="SYSTEM_PROPERTIES_MODE_OVERRIDE" />
  <property name="locations">
    <list>
      <value>classpath:batch.properties</value>
    </list>
  </property>
</bean>

<!-- NCCIP LOADER JOB -->
<job id="load-nccip" xmlns="http://www.springframework.org/schema/batch"
      job-repository="jobRepository-memory">
  <step id="nccipStep">
    <tasklet task-executor="taskExecutor" transaction-manager="transactionManager">
      <chunk reader="nccip-reader" writer="intact-jpa-writer"
              commit-interval="${batch.config.commit-interval}"
              retry-limit="1" skip-policy="skipPolicy">
        <retryable-exception-classes>
          <include class="org.springframework.dao.TransientDataAccessException" />
        </retryable-exception-classes>
        <listeners>
          <listener ref="skipListener" />
        </listeners>
      </chunk>
    </tasklet>
  </step>
  <listeners>
    <listener ref="intactJobExecutionListener" />
  </listeners>
</job>

<!-- RTS LOADER JOB -->
<job id="load-rt" xmlns="http://www.springframework.org/schema/batch"
      job-repository="jobRepository-memory">
  <step id="rtsStep">
    <tasklet task-executor="taskExecutor" transaction-manager="transactionManager">
      <chunk reader="rts-reader" writer="intact-jpa-writer"
              commit-interval="${batch.config.commit-interval}"
              retry-limit="1" skip-policy="skipPolicy">
        <retryable-exception-classes>
          <include class="org.springframework.dao.TransientDataAccessException" />
        </retryable-exception-classes>
      </chunk>
    </tasklet>
  </step>
  <listeners>
    <listener ref="intactJobExecutionListener" />
  </listeners>
</job>
```


# Agenda

- Spring & Spring Batch Overview
- Domain Language of Batch
- Configuring and Running a Job
- **Advanced Features**
- Summary

# Sequential Flow

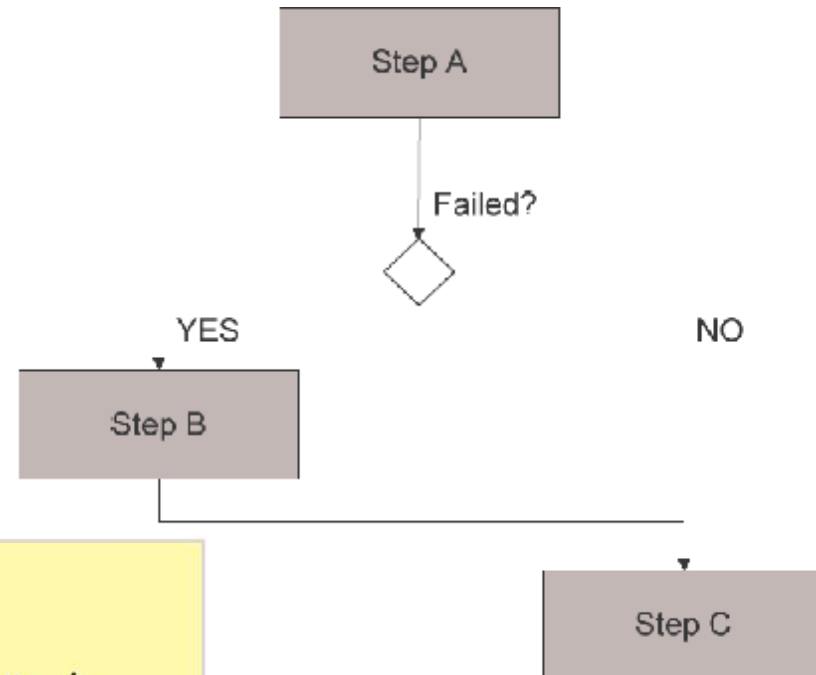
- The simplest flow scenario is a job where all of the steps execute sequentially
- This can be achieved using the 'next' attribute of the step element

```
<job id="job">  
  <step id="stepA" next="stepB" />  
  <step id="stepB" next="stepC" />  
  <step id="stepC" />  
</job>
```



## Conditional Flow

- In order to handle more complex scenarios, Spring Batch allows transition elements to be defined within the step element




```
<job id="job">
  <step id="stepA">
    <next on="FAILED" to="stepB" />
    <next on="*" to="stepC" />
  </step>
  <step id="stepB" next="stepC" />
  <step id="stepC" />
</job>
```

# Split Flow

- Spring Batch also allows for a job to be configured with parallel flows using the 'split' element

```
<split id="stepB" next="stepC">  
  <flow>  
    <step id="stepB11" next="stepB11"/>  
    <step id="stepB12"/>  
  </flow>  
  <flow>  
    <step id="stepB21"/>  
  </flow>  
</split>  
<step id="stepC"/>
```



# Skip & Retry & Fatal Logic

- Based on catching and processing Java exceptions
- Skip Logic
  - There are scenarios where errors encountered should not result in Step failure, but should be skipped instead
  - Avoid step fail when unique constraints violations, ...
- Retry Logic
  - Ability to process one item more than once in order to avoid Step failure
  - Can handle optimistic locking exceptions, ...
- Fatal Exceptions
  - Causes immediate failure when specified exception caught
  - Input file not found, ...

# Intercepting Job Execution

- You might need to perform some functionality at certain events during the execution of a Step
- Based on listeners

```
<step id="step1">  
  <tasklet reader="reader" writer="writer"  
    commit-interval="10"/>  
  <listeners>  
    <listener ref="stepListener"/>  
  </listeners>  
</step>
```

- StepExecutionListener, ChunkListener, ItemReadListener, ItemProcessListener, ItemWriteListener, SkipListener



# Available Item Readers

Item Reader	Description
AbstractItemCountingItemStreamItemReader	Abstract base class that provides basic restart capabilities by counting the number of items returned from an <code>ItemReader</code> .
ListItemReader	Provides the items from a list, one at a time
ItemReaderAdapter	Adapts any class to the <code>ItemReader</code> interface.
AggregateItemReader	An <code>ItemReader</code> that delivers a list as its item, storing up objects from the injected <code>ItemReader</code> until they are ready to be packed out as a collection. This <code>ItemReader</code> should mark the beginning and end of records with the constant values in <code>FieldSetMapper</code> <code>AggregateItemReader#BEGIN_RECORD</code> and <code>AggregateItemReader#END_RECORD</code>
FlatFileItemReader	Reads from a flat file. Includes <code>ItemStream</code> and <code>Skippable</code> functionality. See section on Read from a File
StaxEventItemReader	Reads via StAX. See HOWTO - Read from a File
JdbcCursorItemReader	Reads from a database cursor via JDBC. See HOWTO - Read from a Database
HibernateCursorItemReader	Reads from a cursor based on an HQL query. See section on Reading from a Database
IbatisPagingItemReader	Reads via iBATIS based on a query. Pages through the rows so that large datasets can be read without running out of memory. See HOWTO - Read from a Database
JmsItemReader	Given a Spring <code>JmsOperations</code> object and a JMS Destination or destination name to send errors, provides items received through the injected <code>JmsOperations</code> <code>receive()</code> method
JpaPagingItemReader	Given a JPQL statement, pages through the rows, such that large datasets can be read without running out of memory
JdbcPagingItemReader	Given a SQL statement, pages through the rows, such that large datasets can be read without running out of memory

# Available Item Writers

Item Writer	Description
AbstractItemStreamItemWriter	Abstract base class that combines the <code>ItemStream</code> and <code>ItemWriter</code> interfaces.
CompositeItemWriter	Passes an item to the process method of each in an injected <b>List</b> of <b>ItemWriter</b> objects
ItemWriterAdapter	Adapts any class to the <code>ItemWriter</code> interface.
PropertyExtractingDelegatingItemWriter	Extends <code>AbstractMethodInvokingDelegator</code> creating arguments on the fly. Arguments are created by retrieving the values from the fields in the item to be processed (via a <code>SpringBeanWrapper</code> ) based on an injected array of field name
FlatFileItemWriter	Writes to a flat file. Includes <code>ItemStream</code> and <code>Skippable</code> functionality. See section on Writing to a File
HibernateItemWriter	This item writer is hibernate session aware and handles some transaction-related work that a non-"hibernate aware" item writer would not need to know about and then delegates to another item writer to do the actual writing.
JdbcBatchItemWriter	Uses batching features from a <code>PreparedStatement</code> , if available, and can take rudimentary steps to locate a failure during a <code>flush</code> .
JpaItemWriter	This item writer is JPA <code>EntityManager</code> aware and handles some transaction-related work that a non-"jpa aware" <code>ItemWriter</code> would not need to know about and then delegates to another writer to do the actual writing.
StaxEventItemWriter	Uses an <b>ObjectToXmlSerializer</b> implementation to convert each item to XML and then writes it to an XML file using StAX.

# Agenda

- Spring & Spring Batch Overview
- Domain Language of Batch
- Configuring and Running a Job
- Advanced Features
- **Summary**

## Summary

- Lack of a standard enterprise batch architecture is resulting in higher costs associated with the quality and delivery of solutions.
- Spring Batch provides a highly scalable, easy-to-use, customizable, industry-accepted batch framework
- Spring patterns and practices have been leveraged allowing developers to focus on business logic, while enterprise architects can customize and extend architecture concerns

# Thank you!

– Questions ?

